

SQLite

CS582 Project Phase 2

...

Alex Robic and Matt Bundas

SQLite - An Overview

- SQL-type RDBMS, known for being minimalist, concise, local, yet powerful

”SQLite is a C-language library that implements a small, fast, self-contained, high-reliability, full-featured, SQL database engine. SQLite is the most used database engine in the world. SQLite is built into all mobile phones and most computers and comes bundled inside countless other applications that people use every day.” - sqlite.org

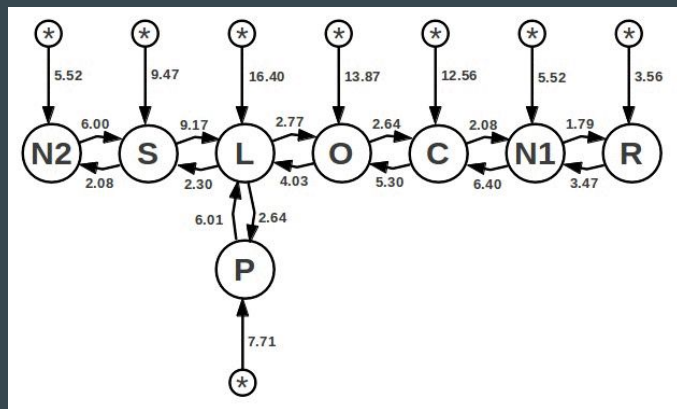
- Used when you are dealing with local, non-enormous, single-writer data

- Not open source, but in the public domain



SQLite Optimization - Finding the Fastest Plan

- SQLite constructs graphs whose nodes are tables, and edges are operations such as joins, with estimated cost along them.
- Original SQLite (pre 2014) used a Nearest Neighbor heuristic to find best plan.
- Modern SQLite uses “N Nearest Neighbor” (N3).
- Typically uses $N = 1$ for simple queries, $N = 5$ for two-way joins and $N = 10$ for larger joins.
- In sample graph, 1NN finds path R-N1-N2-S-C-O-L-P with cost 36.92. $N=8$ finds path with cost 29.78. Actual optimal solution is path with cost 27.38. These costs are logarithmic
- For queries which make use of many OR terms in WHERE clause, SQLite may do a full-scan.



<https://www.sqlite.org/queryplanner-ng.html>

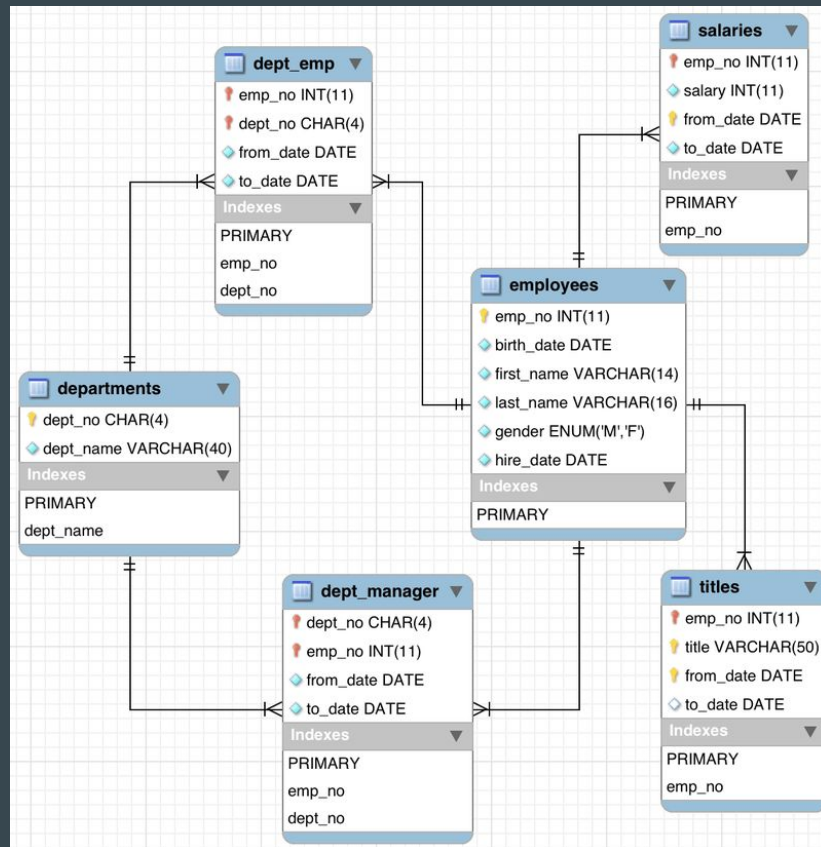
A sample cost graph used in determining query execution plan

SQLite Optimization - Other Ideas

- Stores “rowid”, which is a sorted integer index. This is the value SQLite uses instead of primary key, unless primary key is already an integer. Allows primary key to be NULL.
- Optimizes BETWEEN statements, such as $e1$ between $e2$ and $e3$ into virtual tables. Where it evaluates $e1 \geq e2$ AND $e1 \leq e3$. Behind the scenes it allows SQLite to only load $e1$ once.
- Optimizes OR statements by rewriting them as a single IN statement, or as a series of joins. Helps to optimize usage of indexes.
- When more than one index can be used, optimizer picks option with least amount of work.
- Many more clever tricks with regards to WHERE clauses, covering indices, subquery flattening, co-routines, automatic indexing, min/max optimization.
- Does seem to temporarily store tables created during queries.

Database Used

- An artificially generated database representing a company's departments, department managers, employees and their titles and salaries.
- Provided by MySQL as a large sample database.
- Contains 6 tables, ~4 million entries, about 200 MB.
- https://github.com/datacharmer/test_db



Methods Used

- Testing done on common remote server
- MySQL database constructed using “SOURCE employees.sql”
- SQLite database constructed using “.read employees.sql” on personal computer to create .db file, then transferred to remote server
- Ran each test 8 times, calculated the average.

MySQL

```
mysql> SELECT avg(salary) as Avg_Salary
->
-> FROM salaries where to_date = '2002-08-01';
+-----+
| Avg_Salary |
+-----+
| 71281.1487 |
+-----+
1 row in set (0.26 sec)

mysql> show profiles;
+-----+-----+-----+
| Query_ID | Duration | Query |
+-----+-----+-----+
| 1 | 0.25984350 | SELECT avg(salary) as Avg_Salary |
+-----+-----+-----+
FROM salaries where to_date = '2002-08-01' |
| 2 | 0.26099550 | SELECT avg(salary) as Avg_Salary |
+-----+-----+-----+
FROM salaries where to_date = '2002-08-01' |
| 3 | 0.26106920 | SELECT avg(salary) as Avg_Salary |
+-----+-----+-----+
FROM salaries where to_date = '2002-08-01' |
+-----+-----+-----+
```

SQLite

```
sqlite> .timer on
sqlite> SELECT t.title as Title, avg(s.salary) as Avg_Salary
...
...> FROM salaries as s, employees as e, titles as t
...
...> WHERE e.emp_no = s.emp_no and
...> e.hire_date = s.from_date and
...
...> GROUP BY t.title
...
...> ORDER BY avg(s.salary) desc;
Senior Staff|58572.9697378224
Staff|58502.8904992846
Manager|51531.0416666667
Engineer|48537.7355020257
Assistant Engineer|48530.9304805793
Senior Engineer|48502.1024057054
Technique Leader|48437.3035034772
Run Time: real 0.514 user 0.390625 sys 0.125000
```

Overview Of Queries

- Adapted the themes of the sample queries to the employees database.
- Include joins, sorts, aggregations, filtering, ranging.
- Provided indexes on attributes used in queries which aren't there by default.
- Times range from 4s to 0.0033s.

Query 1

Average Starting Salary of Each Title

```
SELECT t.title as Title, avg(s.salary) as Avg_Salary  
  
FROM salaries as s, employees as e, titles as t  
  
WHERE e.emp_no = s.emp_no and  
      e.hire_date = s.from_date and  
      t.emp_no = e.emp_no  
  
GROUP BY t.title  
ORDER BY avg(s.salary) desc
```

No Additional Indexes

Query Execution Plan No Indexes:

- 1) Scan titles using automatic primary index
- 2) Search employees using automatic index on rowid
- 3) Search salaries using automatic index on emp_no and from_date
- 4) Do group by with B-Tree
- 5) Perform aggregation
- 6) Do order by with B-Tree

With Additional Indexes

Indexes Added:

from_date on salaries
hire_date on employees

Query Execution Plan No Indexes:

- 1) Scan titles using automatic primary index
- 2) Search employees using rowid
- 3) Search salaries using automatic index on emp_no and from_date
- 4) Do group by with B-Tree
- 5) Perform aggregation
- 6) Do order by with B-Tree

Query 1 - Results


title	Avg_Salary
-----	-----
Senior Staff	58572.9697378224
Staff	58502.8904992846
Manager	51531.0416666667
Engineer	48537.7355020257
Assistant Engineer	48530.9304805793
Senior Engineer	48502.1024057054
Technique Leader	48437.3035034772

SQLite No Index - 0.83s

SQLite With Index - 0.85s

MySQL No Index - 4.06s

MySQL With Index - 3.32s



```
SELECT t.title as Title, avg(s.salary) as Avg_Salary
```

```
FROM salaries as s, employees as e, titles as t
```

```
WHERE e.emp_no = s.emp_no and e.hire_date = s.from_date and t.emp_no = e.emp_no
```

```
GROUP BY t.title
```

```
ORDER BY avg(s.salary) desc
```

Query 2

Number of Hires Made Under Each Manager

```
SELECT dm.emp_no as Manager, count(de.emp_no) as Hires

FROM dept_manager as dm

JOIN departments as d ON dm.dept_no = d.dept_no
JOIN dept_emp as de on de.dept_no = d.dept_no

WHERE de.from_date < dm.to_date and
de.from_date > dm.from_date

GROUP BY dm.emp_no

ORDER BY count(de.emp_no) desc;
```

No Additional Indexes

Query Execution Plan No Indexes:

1. Scan dept_manager with automatic primary index
2. Search department with automatic index on dept_no
3. Search dept_emp with automatic on dept_no
4. Filter data based on condition
5. Use B-Tree for Group By
6. Perform Aggregation
7. Use B-Tree for Order By

With Additional Indexes

Indexes Added:

from_date and to_date on dept_manager
from_date and to_date on dept_emp

Query Execution Plan With Indexes:

1. Scan dept_manager with automatic primary index
2. Search department with automatic index on dept_no
3. Search dept_emp with automatic on dept_no
4. Filter data based on condition
5. Use B-Tree for Group By
6. Perform Aggregation
7. Use B-Tree for Order By

Query 2 - Results


Manager	Hires
110567	45360
110511	40322
111133	31506
111035	20725
110386	19517
110420	19513
110344	18069
110303	16342
111534	13702
110039	12451
110114	11808

SQLite No Index - 1.24s

SQLite With Index - 1.19s

MySQL No Index - 2.99s

MySQL With Index - 2.73s



```
SELECT dm.emp_no as Manager, count(de.emp_no) as Hires
```

```
FROM dept_manager as dm
```

```
JOIN departments as d ON dm.dept_no = d.dept_no JOIN dept_emp as de on de.dept_no = d.dept_no
```

```
WHERE de.from_date < dm.to_date and de.from_date > dm.from_date
```

```
GROUP BY dm.emp_no ORDER BY count(de.emp_no) desc
```

Query 3

Average Salary of Employees Named Manu

```
SELECT avg(s.salary) as Avg_Salary  
  
FROM employees as e, salaries as s  
  
WHERE e.first_name = 'Manu' and e.emp_no = s.emp_no
```

No Additional Indexes

Query Execution Plan No Indexes:

- 1) Scan Employees for “Manu”
- 2) Search Salaries With Primary Index on emp_no
- 3) Perform Aggregation

With Additional Indexes

Indexes Added:

first_name on employees

Query Execution Plan With Indexes:

- 1) Search Employees with Index on first_name
- 2) Search Salaries With Primary Index on emp_no
- 3) Perform Aggregation

Query 3 - Results

Avg_Salary


65074.8788492707

SQLite No Index - 0.0725s

SQLite With Index - 0.0039s

MySQL No Index - 0.175s

MySQL With Index - 0.049s



```
SELECT avg(s.salary) as Avg_Salary
```

```
FROM employees as e, salaries as s
```

```
WHERE e.first_name = 'Manu' and e.emp_no = s.emp_no
```

Query 4

Information of Employees With Salary Between 100,000 and 150,000

```
SELECT DISTINCT e.emp_no, e.first_name, e.last_name  
FROM employees as e  
JOIN salaries as s on e.emp_no = s.emp_no  
WHERE s.salary >= 100000 and s.salary <= 150000;
```

No Additional Indexes

Query Execution Plan No Indexes:

- 1) Scan employees table
- 2) Search salaries that meet criteria with automatic primary index on emp_no

With Additional Indexes

Index Added:

salary on salaries

Query Execution Plan With Indexes:

- 1) Search salaries that meet criteria with index on salaries
- 2) Search employees using automatic index on rowid

Query 4 - Results


emp_no	first_name	last_name
10066	Kwee	Schusler
10068	Charlene	Brattka
10087	Xinglin	Eugenio
10107	Dung	Baca
10136	Zissis	Pintelas
10150	Zhenbing	Perng
10151	Itzhak	Lichtner
10160	Debatosh	Khasidashvili
10173	Shrikanth	Mahmud
10185	Duro	Sidhu

SQLite No Index - 1.04s

SQLite With Index - 0.92s

MySQL No Index - 2.63s

MySQL With Index - 0.60s



```
SELECT DISTINCT e.emp_no, e.first_name, e.last_name
```

```
FROM employees as e
```

```
JOIN salaries as s on e.emp_no = s.emp_no
```

```
WHERE s.salary >= 100000 and s.salary <= 150000;
```

Query 5

**Average Salary of Salaries Which Expire on
2002-08-01**

```
SELECT avg(salary) as Avg_Salary  
  
FROM salaries where to_date = '2002-08-01';
```

No Additional Indexes

Query Execution Plan No Indexes:

- 1) Scan salaries table for rows that meet criteria
- 2) Perform aggregation

With Additional Indexes

Index Added:

to_date on salaries

Query Execution Plan With Indexes:

- 1) Search salaries with index on to_date
- 2) Perform aggregation

Query 5 - Results

Avg_Salary


71281.1486880466

SQLite No Index - 0.45s

SQLite With Index - 0.0033s

MySQL No Index - 1.09s

MySQL With Index - 0.0057s




SELECT avg(salary) as Avg_Salary

FROM salaries where to_date = '2002-08-01';

Data from all queries

SQLite										SQLite (with indexing)									
Run number	1	2	3	4	5	6	7	8	Average	Run number	1	2	3	4	5	6	7	8	Average
Q1	0.878	0.895	0.878	0.862	0.752	0.871	0.747	0.756	0.829875	Q1	0.684	0.858	0.88	0.778	0.694	0.939	0.858	1.13	0.852625
Q2	1.143	1.094	1.175	1.183	1.273	1.493	1.364	1.186	1.238875	Q2	1.104	1.087	1.287	1.148	1.15	1.352	1.086	1.271	1.185625
Q3	0.061	0.056	0.079	0.073	0.101	0.067	0.066	0.077	0.0725	Q3	0.003	0.005	0.003	0.004	0.005	0.003	0.004	0.004	0.003875
Q4	1.03	1.079	1.031	0.927	1.012	1.394	0.914	0.923	1.03875	Q4	0.865	0.844	0.95	0.857	0.872	0.941	0.939	1.15	0.92725
Q5	0.426	0.418	0.372	0.348	0.507	0.373	0.679	0.466	0.448625	Q5	0.004	0.003	0.003	0.003	0.004	0.004	0.003	0.003	0.003375
MySQL										MySQL (with indexing)									
Run number	1	2	3	4	5	6	7	8	Average	Run number	1	2	3	4	5	6	7	8	Average
Q1	3.52	4.17	4.02	4.19	3.95	4.07	4.53	4.06	4.06375	Q1	3.79	3.09	3.63	3.12	3.42	3.36	2.92	3.23	3.32
Q2	3.72	2.67	2.79	2.74	3.3	2.98	2.9	2.79	2.98625	Q2	2.43	2.71	2.52	2.82	2.86	2.81	3.14	2.54	2.72875
Q3	0.17	0.23	0.18	0.14	0.16	0.18	0.18	0.16	0.175	Q3	0.00384	0.0072	0.357	0.003	0.00825	0.0063197	0.00364	0.004	0.0491562125
Q4	2.48	2.52	2.96	2.48	2.68	2.63	2.43	2.83	2.62625	Q4	0.85	0.52	0.52	0.63	0.52	0.67	0.47	0.59	0.59625
Q5	1.22	1.14	1	1.13	0.97	0.99	0.94	1.32	1.08875	Q5	0.0098	0.005	0.0048	0.0043	0.0047	0.0047	0.00518	0.0072	0.00571

Total Results



Other Speed Comparisons

Time to create database:

SQLite - 16s

MySQL - 19s

Drop From Salaries where Salary > 125,000:

SQLite No Index- 0.417s

SQLite Index on Salaries - 0.002s

MySQL - 0.65s

MySQL With Index on Salaries - 0.03s

Create index salary on salaries:

SQLite - 1.95s

MySQL - 2.5s

SQLite Demos

SQLite Concurrent Writing:

```
sqlite> .read write_2.sql
Run Time: real 0.076 user 0.078125 sys 0.000000
Error: near line 1: database is locked
Run Time: real 0.107 user 0.078125 sys 0.015625
Run Time: real 0.104 user 0.093750 sys 0.000000
Run Time: real 0.105 user 0.078125 sys 0.000000
Run Time: real 0.118 user 0.093750 sys 0.000000
Run Time: real 0.099 user 0.078125 sys 0.000000
Error: near line 124601: database is locked
Run Time: real 0.088 user 0.062500 sys 0.000000
Error: near line 149521: database is locked
Run Time: real 0.116 user 0.093750 sys 0.015625
Run Time: real 0.115 user 0.093750 sys 0.000000
Run Time: real 0.117 user 0.093750 sys 0.000000
Run Time: real 0.129 user 0.093750 sys 0.031250
Run Time: real 0.112 user 0.093750 sys 0.015625
Run Time: real 0.128 user 0.078125 sys 0.000000
Run Time: real 0.115 user 0.078125 sys 0.000000
Run Time: real 0.117 user 0.078125 sys 0.000000
Run Time: real 0.124 user 0.109375 sys 0.000000
Run Time: real 0.114 user 0.078125 sys 0.015625
Run Time: real 0.142 user 0.109375 sys 0.000000
Run Time: real 0.121 user 0.093750 sys 0.015625
Run Time: real 0.107 user 0.062500 sys 0.031250
Run Time: real 0.113 user 0.093750 sys 0.000000
Run Time: real 0.129 user 0.078125 sys 0.000000
Run Time: real 0.106 user 0.078125 sys 0.015625
Run Time: real 0.115 user 0.078125 sys 0.015625
Run Time: real 0.114 user 0.093750 sys 0.000000
```

SQLite Typing - Live

Conclusion

- In the five queries, SQLite was 2-13 times faster than MySQL.
- In other tests, SQLite are closer in performance, but with SQLite edging MySQL out.
- Increased performance comes from lack of overhead, simplistic code, dealing with a single .db file.
- Comes through on its selling point.

Resources

- [1] datacharmer. *Employees test_{db}*. URL: https://github.com/datacharmer/test_db.
- [2] Joseph (Yossi) Gil. *LaTeX 2_ε for Graduate Students*. manuscript. Haifa, Israel, 2002.
- [3] MySQL. *Employees Sample Database*. URL: <https://dev.mysql.com/doc/employee/en/>.
- [4] Skookum. *SQLite: The Database at the Edge of the Network with Dr. Richard Hipp*. 2015. URL: https://www.youtube.com/watch?v=Jib2AmRb_rk.
- [5] SQLite. *Appropriate Uses For SQLite*. URL: <https://sqlite.org/whentouse.html>.
- [6] SQLite. *CREATE INDEX*. URL: https://sqlite.org/lang_createindex.html.
- [7] SQLite. *Database Speed Comparison*. URL: <https://sqlite.org/speed.html>.
- [8] SQLite. *Most Widely Deployed and Used Database Engine*. URL: <https://www.sqlite.org/mostdeployed.html>.
- [9] SQLite. *Rowid Tables*. URL: <https://www.sqlite.org/rowidtable.html>.
- [10] SQLite. *The Next-Generation Query Planner*. URL: <https://www.sqlite.org/queryplanner-ng.html>.
- [11] SQLite. *The SQLite Query Optimizer Overview*. URL: <https://www.sqlite.org/optoverview.html>.
- [12] SQLITETUTORIAL. *SQLite Index*. 2021. URL: <https://www.sqlitetutorial.net/sqlite-index/>.